

The Pi-Calculus

2 March 2021

0. What motivated the search for a process calculus?
1. What did prehistoric models look like?
2. How did the pi-calculus improve upon them?
3. What is The Truth for a process calculus?
4. How do you add discipline to the world?

Pre-History

In the late 70s / early 80s, researchers had a model of computation (λ -calculus), but no model for concurrent communication.

Researchers wanted to reason about processes running independently and in communication with each other.

Many 'process calculi' emerged:

- Actors
- CSP
- Petri Nets
- ...

Many models, little agreement.

1. Milner 1980 - CCS

Calculus of Communicating Systems (CCS).

Two Ideas:

1. A process can be defined based on how an Observer interacts with it. To build complex systems, you can take an observer and model it as a process that can be observed.
2. Processes communicate over channels pathways over which data $\xrightarrow{\text{can be sent}}$ between processes.

→ NOT a new idea.

Intro to CCS

Preliminaries:

channels are defined by a set of names, unique symbols that differ from other values. We use $\alpha, \beta, \gamma, \delta, \dots$ and range over channels with μ .

Denote processes with P, Q, R .

We define processes via their interactions with processes over channels.

Channels are a static property of a process.



Example: checkout at a grocery store.

1. nil

- a closed checkout aisle.



2. receive

- the cashier receives a payment from customer.

$$\alpha(x). P \xrightarrow{\alpha s} P'[x \leftarrow s]$$

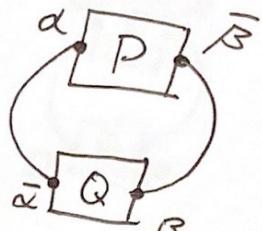
3. send

- the cashier hands you your purchase.

$$\bar{\beta} 10. P \xrightarrow{\beta 10} p'$$

4. Parallel (Par)

- a cashier and customer communicating.

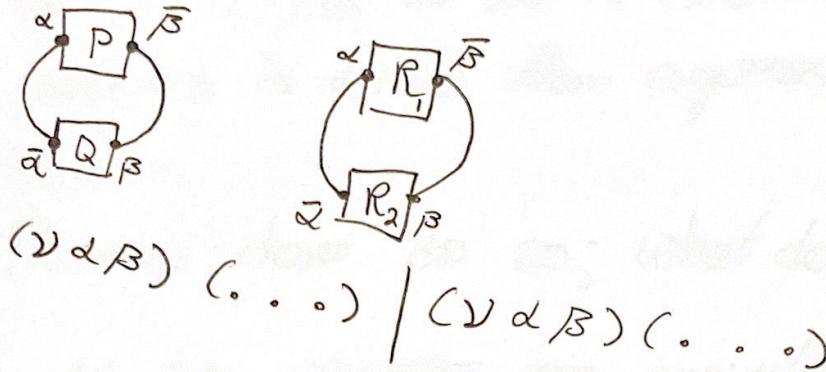


$$\begin{array}{l} (\alpha(x). \bar{\beta} \cdot x. P') \parallel \bar{\delta}. \beta(y). Q' \\ \xrightarrow{\alpha s} (\bar{\beta} \cdot 10. P'[x \leftarrow s] \mid \beta(y). Q') \\ \xrightarrow{\beta 10} (P'[x \leftarrow s] \mid Q'[y \leftarrow 10]) \end{array}$$

Note: From the perspective of the observer, they don't interact with the system. These are sometimes called τ -transitions ($\xrightarrow{\tau}$). This distinction is only useful for theoretical purposes.

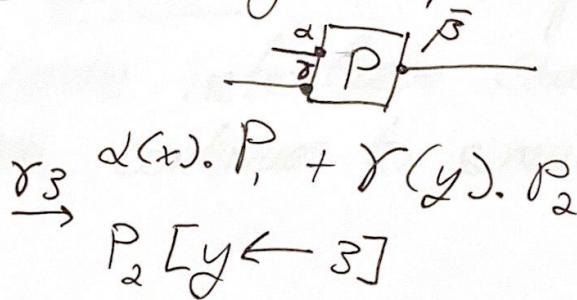
5. Restriction

- cashiers & customers in different aisles.

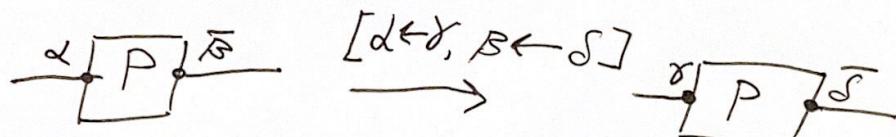


6. Summation

- a cashier might respond in multiple ways.



7. Relabelling



$$\begin{aligned} & (\alpha(x). \bar{\beta} 2^x. P') [\alpha \leftarrow r, \beta \leftarrow s] \\ \rightarrow & (\gamma(x). \bar{s} 2^x. P' [\dots]) \end{aligned}$$

Relabelling is a part of the syntax. Channels are static, and renaming & substitution isn't defined over them.

Problem: where did the Truth go?

In the λ -calculus, the Truth is Observational Equivalence.
Its meaning is derived from expressions reducing to values.

Processes don't do so; what do we do?

Insight: two processes are equivalent if they simulate each other.

A process simulates another process if it has the same interface (channels) and, after a transition, continues to simulate the other process.

This is called a bisimulation/bisimilarity relation.

Reflections on CCS

Milner was unhappy for two reasons:

1. Channels are static. In the real world communication links change dynamically with time.
2. Could a smaller calculus be developed?

Intermediate Step - Ulféberg & Nielsen, 1986

Solution: allow processes to pass channels over channels.

Extended CCS (or ECCS).

Two disjoint sets of variables:

- value variables: $x, y, z \dots$
 - name/channel variables: a, b, c

Example: Walkie-Talkies

Two processes would like to communicate over a private talk channel.

$$\begin{array}{c}
 \frac{\beta \quad \beta}{\bar{x} \text{ talk}. \text{ talk}(x). P' / \alpha(\text{talk}). \bar{s}. Q'} \\
 \xrightarrow{\alpha \beta} \\
 \beta(x). P' \quad | \quad \bar{\beta} s. Q' [\text{talk} \leftarrow \beta] \\
 \xrightarrow{\beta s} \\
 P'[x \leftarrow s] / Q'[\text{talk} \leftarrow \beta]
 \end{array}$$

ECCS Reflection

The capacity for the links between processes to change dynamically is called 'mobility.'
↳ channels no longer unidirectional.

This also meant shrinking the calculus. We no longer need relabelling, since we've defined substitution and renaming over channel variables.

2. Milner, Parrow, Walker 1992 - π -calculus

Small improvement: distinguishing between channel
and value variables is unnecessary to
achieve the same expressiveness.

The relatively small change of removing the
distinction is what produces the
 π -calculus.

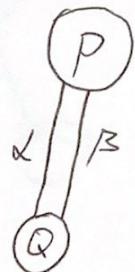
Characteristic Examples

1. Passing Channels

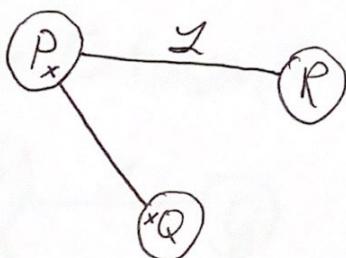
$$\overline{\alpha \beta} \cdot \beta(x). P' \mid \alpha(\text{chan}). \text{chan } s. Q'$$

$$\xrightarrow{\alpha \beta} \beta(x). P' \mid \overline{\beta} s. Q'[\text{chan} \leftarrow \beta]$$

$$\xrightarrow{\beta s} P'[x \leftarrow s] \mid Q'[\text{chan} \leftarrow \beta]$$



2. Restricting Channels



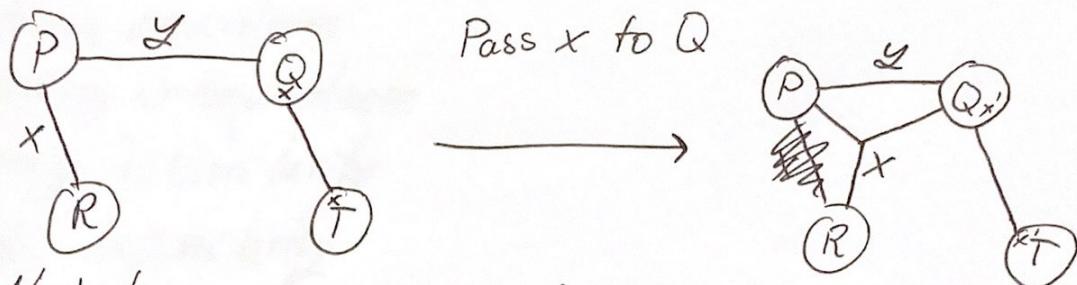
$$\xrightarrow{x5} (\nu x) (\overline{x} s. P' \mid x(y). Q') \mid R$$

$$\rightarrow (\nu x) (P' \mid Q'[y \leftarrow s]) \mid R$$

A well-defined process (like a λ -calculus program)
has all channels explicitly introduced by restriction.

Characteristic Examples 2: Breaking the Scope Barrier

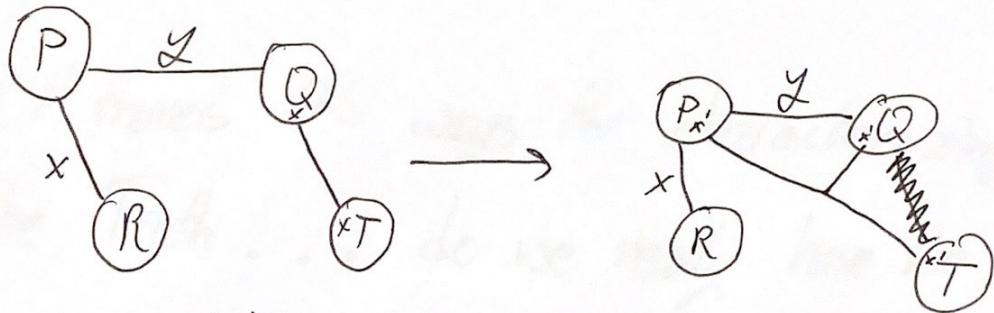
3. Scope Intrusion



Need to rename restricted channels.

$$\begin{aligned} & \bar{y}x. P' / R / (\nu x) (\bar{y}z. Q' / T) \\ \xrightarrow{\bar{y}x} & P' / R / (\nu x) (Q'[x \leftarrow x'][z \leftarrow z] / T[x \leftarrow x']) \end{aligned}$$

4. Scope Extrusion



$$\bar{y}z. P' / R / (\nu x) (\bar{y}x. Q' / T)$$

$$\xrightarrow{\bar{y}x'} P'[\bar{z} \leftarrow x'] / R / (\nu x') (Q'[x \leftarrow x'] / T[x \leftarrow x'])$$

Final Note on The Truth:

Milner et al. introduce Five bisimilarity relations.

- strong bisimilarity
- strong equivalence
- Strong D-Equivalence
- early bisimilarity
- late bisimilarity

And there's more where that came from:

- barbed congruence
- strong barbed congruence

• • •

If there's 25 ways for characterizing
the Truth . . . do we really have the Truth?

3. Milner 1993 - Adding Discipline

The π -calculus is powerful, but chaotic.

Example: the 'molecular action.'

$u, v \in \text{values}$

$x \in \text{channels}$

$$\bar{x} u. \bar{x} v. P' | x(y). x(z). Q' | x(y). x(z). R'$$

Values are NOT broadcasted over channels.

$$\xrightarrow{xu} \bar{x} v. P' | x(z). Q'[y \leftarrow u] | x(y). x(z). R'$$

$$\xrightarrow{xv} P' | x(z). Q'[y \leftarrow u] | x(z). R'[z \leftarrow v]$$

The values can get split between processes
 Q and R .

How can we ensure they're received by the
same process?

Solution: send a channel first.

$x, w \in \text{channels}$.

$$\begin{array}{l} \overline{x} w. \overline{w} u. \overline{w} v. P' / x(w). \overline{w}(u). \overline{w}(v). Q' / x(w). \overline{w}(u). w(v). R' \\ \xrightarrow{w} \overline{w} u. \overline{w} v. P' / w(u). w(v). Q' / x(w). w(u). w(v). R' \\ \xrightarrow{w} \overline{w} v. P' / w(v). Q' / x(w). w(u). w(v). R' \\ \xrightarrow{wv} P' / Q' / x(w). w(u). w(v). R' \end{array}$$

Expressive & powerful, but intricate.

Intermediate Step: Polyadic Π -calculus

Instead of sending single values, send and receive tuples of arbitrary arity.

$$P = \bar{x} \langle u, v \rangle . P'$$

$$Q = x(u, v) . Q'$$

$$R = x(u, v) . R'$$

- Good: can pass multiple values at a time.
- Bad: Arity errors.

How do we add discipline?

A Sorting Discipline

Sorts are analogous to types in λ -calculus.

Each channel is associated with a sort. A sort is mapped to an 'Object sort,' or a list of sorts.

Intuition: we express the arities of each channel, and the arities of the values passed over it.

$$S = \{ \text{SENDER} : [\text{LEFT}, \text{RIGHT}], \text{LEFT} : [], \text{RIGHT} : [] \}$$

$$(2) x : \text{SENDER}, u : \text{LEFT}, v : \text{RIGHT}) (. . .)$$

Note: We need names because our types are recursive.

Inference Rules:

We show that processes are well-typed. A process must be well-defined.

A closed process has sort $[]$.

A process with a free channel has the sort of that channel.

$$\frac{\Gamma \vdash x : S \quad \Gamma \vdash P : \text{ob}(S)}{\Gamma \vdash x.P : []} \quad [\text{T-IN}]$$

$$\frac{\Gamma \vdash x : S \quad \Gamma \vdash P : \text{ob}(S)}{\Gamma \vdash \bar{x}.P : []} \quad [\text{T-OUT}]$$

$$\frac{\Gamma \vdash P : [] \quad \Gamma \vdash Q : []}{\Gamma \vdash (P|Q) : []} \quad [\text{T-PAR}]$$

$$\frac{\Gamma, x : S \vdash P : []}{\Gamma \vdash (\lambda x : S)P : []} \quad [\text{T-RESTR}]$$

Reflection

The sorting discipline guarantees no arity runtime errors.

Is this what we want?

This discipline is unsatisfactory. We haven't learned how to use channels.

Example: what if type systems for modern programming languages only let the programmer specify the arities of functions?

Arity-checking is insufficient.

Pierce & Sangorgi 1994 - I/O Types

Motivation: the office printer.
job, paper ∈ channels

Printer = job(x). paper(x). Printer

If somebody else is reading on the job channel,
they can steal the data somebody else wanted to print.

Solution: specify what direction the channel can
be used for a process.

Preliminaries:

Polarities $\{1, 0\}$

Polarities = $\{1, 0\}$

$r = \{1\}$

$w = \{0\}$

$b = \{1, 0\}$

e.g. $r[T_1, T_2, \dots, T_n]$

↳ a sort for a channel of arity n that can only be read.

Printer example:

JobS001S01

(\triangleright job: $r[S]$, paper: $w[S]$) ($\text{job}(x)$. $\overline{\text{paper}}(x)$. Printer)

A client Q will receive the job channel & specific interface:

$Q = \text{recv}(\text{job}: w[S]). \overline{\text{job}} S. Q'$

The Sub-typing relation:

We need a way to show that a sort obeys the interface of a channel

$$b <: r$$

$$b <: w$$

$<:$ A sort S is a subsort of T iff:

1. if $T = b[T_1 \dots T_n]$ then $S = b[S_1 \dots S_n]$ and both $S_i <: T_i$ and $T_i <: S_i$.

2. if $T = r[T_1 \dots T_n]$ then $S = r[S_1 \dots S_n]$, $r <: r$, and for all $1 \leq i \leq n$, $S_i <: T_i$.

3. if $T = w[T_1 \dots T_n]$ then $S = w[S_1 \dots S_n]$, $w <: w$, and for all $1 \leq i \leq n$, $T_i <: S_i$.

Example Rule: (Σ is the subsort context) \dagger

$$I <: r \quad \text{for each } i, \Sigma \vdash S_i <: T_i$$

$$\sum \vdash I[S_1 \dots S_n] <: r[T_1 \dots T_n]$$

Typing processes with I/O Types:

Two new rules:
modified

$$\frac{\vdash \Gamma(a) \leqslant r[s_1, \dots, s_n] \quad \Gamma, b_1 : s_1, \dots, b_n : s_n \vdash P : ()}{\Gamma \vdash a(b_1 : s_1, \dots, b_n : s_n). P : ()} \quad [\Gamma\text{-IN}]$$

$$\frac{\vdash \Gamma(a) \leqslant w[\Gamma(b_1) \dots \Gamma(b_n)] \quad \Gamma \vdash P : ()}{\Gamma \vdash \bar{a} \langle b_1, \dots, b_n \rangle. P : ()}$$

Kobayashi, Pierce, Turner 1998 - Linear Pi-Calculus

Linear Types for Channels

Consider a web server.

- it receives any number of requests, and a channel on which to send a response.
- it always sends one response.

The interface / specification of our protocol doesn't just include a channel's direction — it also includes how frequently we can use a channel.

How do we introduce this to the Π -calculus?

Introduction to Linear Types

Linear types let us specify the usage of resources. We assign to a resource a 'multiplicity' and that resource must be 'used' that number of times.

More specifically, narrowing and contraction are disallowed ~~inference~~ with a linear type.

To express this, inference rules are expressed as the sum of two separate environments, to highlight that the resource is used correctly between the two.

$$\frac{\Gamma_1 + v \in V \quad \Gamma_2 + w \in W}{\Gamma_1 + \Gamma_2 + \langle v, w \rangle \in V \times W}$$

Preliminary:

m is a metavariable ranging over multiplicities ($1, \omega$).

e.g. $r[T_1 \dots T_n]^1$

Combination of types:

$$I_1[T_1 \dots T_n]^\omega + I_2[\cancel{T}_1 \dots \cancel{T}_n]^\omega = (I_1 \cup I_2)[T_1 \dots T_n]^\omega$$

$$I_1[T_1 \dots T_n]^1 + I_2[T_1 \dots T_n]^1 = (I_1 \cup I_2)[T_1 \dots T_n]^1$$

if $I_1 \cap I_2 = \emptyset$

The capabilities of our channels must be split: a linear channel can be read from 1 process and written from 1 process, but no more.

Typing Rules:

$$\frac{\Gamma_1 \vdash P \quad \Gamma_2 \vdash Q}{\Gamma_1 + \Gamma_2 \vdash P/Q}$$

[E-PAR]

$$\frac{\Gamma \text{ unlimited}}{\Gamma + x : w[\tau_1 \dots \tau_n]^m + y : \tau_1 \dots \tau_n \vdash \bar{x}[y]}$$

[E-OUT]

$$\frac{\Gamma, z : \tau_1 \dots \tau_n \vdash P \quad \Gamma \text{ unlimited}}{\Gamma + x : r[\tau_1 \dots \tau_n]^m + x(z) . P}$$

[E-IN]

$$\frac{\Gamma, x : I[\tau_1 \dots \tau_n]^m \vdash P}{\Gamma \vdash (\lambda x : I[\tau_1 \dots \tau_n]^m) P}$$

[E-NEW]

$$\frac{\Gamma \vdash P_1 \quad \Gamma \vdash P_2}{\Gamma + b : \text{Bool} \vdash ,^P b \text{ then } P_1 \text{ else } P_2}$$

[E-IF]

Summary

- The key component to the process calculus is mobility, the ability for the communication between processes to evolve dynamically.
- Our notion of the truth has changed — processes simulate one another, and can thus be equated ... somehow.
- We can specify our processes across multiple dimensions, particularly by specifying the interface of a channel.

In Retrospect

The π -calculus has fallen short in two ways:

1. There is no notion of the single π -calculus. For every paper on the subject, you have a different calculus, with varying central operators, and different notions of bisimulation.
2. Nobody programs in the π -calculus. Channels are a common component of process calculi. The π -calculus describes concurrent communication via symbol-pushing, and that doesn't line up with real concurrent programs.